

HEINRICH-HEINE-UNIVERSITÄT DÜSSELDORF

**Bachelor Thesis**

**Implementation and evaluation of the  
reinforcement learning methods Self-Play and  
League Training in the MicroRTS  
environment**

**Niklas Glaser**

|                |                                             |
|----------------|---------------------------------------------|
| Degree:        | B.Sc. Computer Science                      |
| Supervisor:    | Professor Dr. Paul Swoboda                  |
| Co-Supervisor: | Dr. Peter Arndt                             |
| Faculty:       | Faculty of Mathematics and Natural Sciences |


01.04.2026

## Abstract

This thesis improves an existing agent for the real-time strategy (RTS) game MicroRTS. RTS games are considered a challenging domain for Deep Reinforcement Learning (DRL), since agents must deal with large combinatorial action spaces, simultaneous decisions and delayed and often sparse rewards. We implement League Training (LT) in combination with Proximal Policy Optimization (PPO) to address a common problem in DRL training setups such as pure PPO: after sufficient training, the agent reaches a point where it cannot be trained further efficiently with PPO alone, because there are not enough strong opponents to train against or because the available opponents are not diverse enough to develop a robust policy. An obvious and simple countermeasure is Self-Play (SP). However, in practice it can be unstable and lead to overfitting or cyclical strategies, which can reduce performance against external opponents.

LT stabilizes SP by building a population of agents across different training stages and roles. The training agent is deliberately matched against diverse opponents, including specialized agents that are trained to exploit weaknesses of the current policy. We evaluate the resulting agents in matches against various opponent types and compare LT with our initial agent. Our results suggest that LT increases the robustness and generalization of the agent against different opponents, but that its effectiveness depends strongly on the strategic diversity of the initial agents and training opponents. They also suggest that using an agent pre-trained with Behavioral Cloning (BC) as the starting point for LT may be suboptimal when the BC phase is based on only a small and strategically limited set of experts.

Overall, especially after fine-tuning, the evaluation against all opponents shows a substantial improvement of LT over the Base Agent with a BC and PPO approach.



Comparison of total win rates for the Base Agent, the Main Agent after LT and the Main Agent after fine-tuning against all our evaluation opponents.

## Contents

|                                              |           |
|----------------------------------------------|-----------|
| <b>1. Introduction.</b>                      | <b>1</b>  |
| <b>2. Related Work.</b>                      | <b>2</b>  |
| 2.0.1 Game-Playing Agents in MicroRTS        | 2         |
| 2.0.2 MicroRTS-Py (Gym- $\mu$ RTS)           | 3         |
| 2.0.3 RAISocketAI                            | 3         |
| 2.0.4 BeClone                                | 4         |
| 2.0.5 AlphaStar                              | 4         |
| <b>3. Methods.</b>                           | <b>4</b>  |
| 3.1. MicroRTS-environment                    | 4         |
| 3.2. Self-Play (SP)                          | 6         |
| 3.3. Fictitious Self-Play (FSP)              | 6         |
| 3.4. Prioritized Fictitious Self-Play (PFSP) | 7         |
| 3.5. League Training (LT)                    | 8         |
| 3.6. Training Pipeline.                      | 9         |
| 3.7. Additional Modifications for Diversity  | 10        |
| 3.8. Delta Score.                            | 10        |
| <b>4. Experiments.</b>                       | <b>10</b> |
| 4.1. Hyperparameters                         | 11        |
| 4.2. Evaluation                              | 11        |
| 4.3. Results                                 | 11        |
| 4.4. Discussion                              | 13        |
| 4.5. Future Work.                            | 15        |
| <b>5. Conclusion.</b>                        | <b>16</b> |
| <b>Erklärung</b>                             | <b>25</b> |

## 1. Introduction.

Reinforcement Learning (RL) has established itself in recent years as a powerful approach for developing agents in complex environments. Game environments such as RTS games pose a particular challenge in this regard, with large action spaces, long-term planning and simultaneous decisions. However, they also provide clearly defined state spaces, reward structures and repeatable experiments, which makes them an ideal and challenging test environment for researching RL methods.

**MicroRTS.** MicroRTS is a simplified version of RTS games that was developed specifically for research [Huang et al., 2021]. It is discrete and reduces the visual and mechanical complexity of large commercial RTS games, such as StarCraft II. At the same time, it retains the core difficulties, for example resource management, unit production, combat mechanics and tactical positioning. MicroRTS is therefore suitable for the

systematic investigation of learning methods. MicroRTS also contains numerous training challenges, such as large combinatorial action spaces, delayed rewards and multi-agent interactions. It also requires robust strategies against different opponents. For instance, the agent must decide early whether it wants to play aggressively or defensively against the current opponent, which affects the overall game strategy and requires long-term planning.




Figure 1. Screenshot of the initial state of the basesWorkers16x16A map for MicroRTS.

**Bots.** The built-in bots PassiveAI, WorkerRushAI, LightRushAI, RandomAI and RandomBiasedAI, as well as CoacAI, are scripted opponents based on heuristic rules and are often used as opponents for training and evaluation in MicroRTS [Huang et al., 2021]. Mayari is another strong competition bot used in this thesis. The other bots in our evaluation are MCTS-based bots, with a search-based approach to select actions [Huang et al., 2021]. Our aim is to train an agent that can achieve high win rates against these bots, without giving the agent a specific strategy to use against them.

**Self-Play-Based Approaches.** A central problem in RL is generalization across different opponents and strategies: agents often learn strategies that are strongly tailored to the training opponent or a specific game situation. This leads to overfitting and limited robustness. SP, Fictitious Self-Play (FSP) and LT are approaches that address this problem.

**League Training (LT).** LT further stabilizes training by exposing the agent to a league of diverse opponents instead of a single self-play opponent [Vinyals et al., 2019]. This promotes robustness across a broader set of strategies.

**Thesis Outline.** Despite the theoretical advantages, the practical implementation of LT in complex environments such as MicroRTS is associated with significant challenges, that need to be addressed. These include

- the efficient management and selection of opponents for training the Main Agent and for the specialized agents,
- stabilizing the training process under non-constant opponent distributions,
- dealing with large action spaces and sparse or delayed rewards,
- ensuring opponent diversity through sufficient exploration and preventing overfitting to specific opponents or strategies.

In addition, we will compare and evaluate the results empirically. In particular, we will analyze the strategic diversity and robustness for the LT approach and the effects of the initial agent used for the training.

## 2. Related Work.

### 2.0.1 Game-Playing Agents in MicroRTS

The unpublished bachelor thesis by M. Huft, on which this work builds, develops a competitive MicroRTS agent and evaluates the hybrid training pipeline of BC, BC-FT and PPO used to train the Base Agent [Huft, 2025]. This Base Agent is the initial agent for the SP- and LT-based methods in this thesis, unless otherwise specified.

- **BC, BC-FT:** Learning from scratch with DRL methods such as PPO can be very time-consuming [Martin and Jhala, 2021]. However, this training approach reduces this initial, often inefficient exploration in DRL by quickly learning a competent base strategy through imitating expert behavior [Torabi et al., 2018].
- **PPO:** PPO is a well-known on-policy policy optimization method that clips the policy update size to stabilize policy updates [Schulman et al., 2017]. This method further improves the performance of the BC agent through exploration and adaptation to the environment, allowing the agent to adapt and generalize beyond the expert strategies [Martin and Jhala, 2021].

This two-stage approach, in which a policy pre-trained with BC is refined by a DRL algorithm, has proven effective in reducing training time and improving overall performance compared to a pure DRL approach [Martin and Jhala, 2021].

**Results and Limitations.** The results show a clear performance increase across all three training phases (Table 1): BC mainly learns basic game mechanics, while BC-FT improves generalization against multiple bots. Finally, PPO achieves high win rates against many of the MicroRTS bots [Huft, 2025]. There are still agents in the results against which the Base Agent does not perform well. The thesis also identifies limitations, in particular the focus on a single map and the limited opponent diversity and proposes extensions using SP or league mechanisms. Exactly this proposal is taken up by the present work by extending the existing Base Agent and the initial code with SP and LT.

Table 1. Win rates (%) for BC, BC-FT and PPO reported in [Huft, 2025]. The Base Agent in this thesis is trained using the same training pipeline, however these results are of a different version of the Base Agent.

| Bot            | BC | BC-FT | PPO |
|----------------|----|-------|-----|
| WorkerRushAI   | 3  | 50    | 99  |
| PassiveAI      | 98 | 100   | 100 |
| LightRushAI    | 22 | 93    | 99  |
| CoacAI         | 0  | 8     | 96  |
| Mayari         | 0  | 5     | 95  |
| RandomAI       | 52 | 95    | 99  |
| RandomBiasedAI | 18 | 68    | 86  |
| rojo           | 0  | 34    | 84  |
| mixedBot       | 4  | 44    | 59  |
| izanagi        | 0  | 62    | 61  |
| tiamat         | 11 | 78    | 84  |
| droplet        | 0  | 31    | 59  |
| guidedRojoA3N  | 4  | 42    | 75  |
| naiveMCTS      | 0  | 5     | 3   |

## 2.0.2 MicroRTS-Py (Gym- $\mu$ RTS)

Huang et al. provide an efficient Gym-compatible interface for our environment MicroRTS [Ontañón, 2013] with Gym- $\mu$ RTS/MicroRTS-Py [Huang et al., 2021]. MicroRTS-Py is an established and widely used DRL baseline that enables the investigation of DRL approaches in full real-time strategy game scenarios without requiring extensive technical infrastructure such as high-performance computing clusters [Huang et al., 2021].

**Proximal Policy Optimization (PPO) Training.** The openai/baselines [Dhariwal et al., 2017] implementation of PPO was adapted in order to fit MicroRTS-Py. We use a large part of this implementation in LT [Huang et al., 2021]. It was also previously used in the PPO training phases of the Base Agent [Huft, 2025]. On top of openai/baselines, invalid action masking and suitable action/observation representations (e. g. GridNet [Han et al., 2019]) as well as architectural choices (such as IMPALA-CNN, a deep convolutional neural network structure with residual blocks for processing grid-based state representations in RL) [Espeholt et al., 2018] were added [Huang et al., 2021]. In a modified form, these are still used in the LT code.

**GridNet.** Huang et al. considered two action representations:

- **UAS:** UAS iterates over all units at each step and returns the possible actions for each unit, which substantially reduces the action space.
- **GridNet:** GridNet predicts an action for each cell on the map. This means that GridNet executes its large action space in a single step, but this also simplifies the training [Huang et al., 2021, Han et al., 2019].

Against the same evaluation bots, the PPO implementation with invalid action masking and IMPALA-CNN achieves an average win rate of 91% with UAS and 84% with GridNet. Ultimately, the decision was made to focus on GridNet, since its lower complexity and its higher compatibility with SP (and LT) [Huang et al., 2021]. SP was also evaluated in the work. However, there was a bug in the code that likely affected the results for SP [Goodfriend, 2024].

## 2.0.3 RAISocketAI

In 2023 RAISocketAI was the first DRL agent which won the IEEE MicroRTS Competition. RAISocketAI includes a re-implementation of MicroRTS-Py, which is used in our code. Among other things, data generation via replays is an important basis for BC [Huft, 2025] and a bug in the old implementation was fixed that marked non-owned resources as owned resources when the agent is player 2. According to Goodfriend [Goodfriend, 2024], this likely caused SP in the previous MicroRTS-Py implementation to yield no improvements. The RAISocketAI implementation of FSP achieved an average win rate of 91% against the same MicroRTS bots as in MicroRTS-Py, which is significantly higher than 84% in the previous implementation. However, FSP achieved only a 20% win rate against CoacAI, whereas the best previous implementations were almost always able to defeat CoacAI. With fine-tuning on CoacAI and

Mayari, the average win rate against the MicroRTS bots was increased to 98% [Goodfriend, 2024].

#### 2.0.4 BeClone

BeClone demonstrated the effectiveness of BC with DRL fine-tuning in MicroRTS-Py using a two-phase training strategy for RTS agents: first, an initial policy is learned from experts via BC, and subsequently an RL fine-tuning is performed using Advantage Actor-Critic (A2C) [Martin and Jhala, 2021].

#### 2.0.5 AlphaStar

AlphaStar is a DRL agent for the complex RTS game StarCraft II, which was developed by DeepMind [Vinyals et al., 2019]. It demonstrates that LT can decisively contribute to robustness in complex RTS domains. The agent in AlphaStar was pre-trained with imitation learning from human replays.

**Pre-League Training.** In this thesis, we replace imitation learning with the BC phases and the PPO phase. Similar to AlphaStar our agent was further trained in a league of different types of league agents:

- **Main Agent**
- **Main Exploiters**
- **League Exploiters**
- **Historical Agents** (historical snapshots of the Main Agent and the exploiters)

These agents are added to the league to reduce overfitting and forgetting and to promote diversity.

**League Composition and Training.** As in AlphaStar, the Main Exploiters in this thesis are trained against historical versions of the Main Agent in the league in order to identify and exploit weaknesses. At the same time, League Exploiters train against all Historical Agents. The Main Agent trains with Prioritized Fictitious Self-Play (PFSP) against historical versions of itself, Main Exploiters and League Exploiters. The prioritization of the historical opponents is computed based on the Main Agent’s win rate against the respective Historical Agents. In this thesis, the prioritization formula was adjusted to better fit the MicroRTS environment. An exploiter is added to the league as a Historical Agent only when it shows a high win rate against all opponents or when a predefined step limit is reached [Vinyals et al., 2019].

**Differences from AlphaStar.** Although MicroRTS and StarCraft II substantially differ in complexity and representation, AlphaStar is an established reference design for stabilizing and generalizing SP through league-based matchmaking mechanisms and the inclusion of specialized exploiters [Vinyals et al., 2019]. AlphaStar performs policy updates using an actor-critic method with an off-policy correction mechanism called V-trace [Espeholt et al., 2018, Vinyals et al., 2019]. This approach is better suited for LT than PPO due to its robustness to off-policy data and its typically better scalability in large distributed setups. However, the implementation of PPO in MicroRTS-Py is significantly simpler and faster to realize, which is why PPO is a practical choice for implementing LT [Vinyals et al., 2019]. In addition, AlphaStar uses TD( $\lambda$ ) [Sutton and Barto, 1998] for value-function estimation.

### 3. Methods.

#### 3.1. MicroRTS-environment

MicroRTS is a two-player zero-sum game that is played on a two-dimensional grid (the map). Different units can be located on each cell of this map and players can execute actions to move/train/attack units and to collect resources. The different unit types are:

- **Base:**  The base can store resources and train workers. It is very important, since it is the only way to store resources. (cost (in resources): 10, hit points (hp): 10)
- **Resources:**  There is only one type of resource. It can be stored by workers in the base (Number of resources in a cell: 25).
- **Worker:**  workers can collect resources and build barracks (cost: 1, hp: 1, attack strength (at): 1).
- **Barracks:**  barracks can train light, heavy and ranged units (cost: 5, hp: 4).
- **Light:**  light units can move quickly, are trained quickly and cost little (cost: 2, hp: 4, at: 2).
- **Heavy:**  heavy units can move slowly, are trained slowly and cost a lot, but have high at and hp (cost: 3, hp: 8, at: 4).
- **Ranged:**  ranged units have an attack range of 3 cells (cost: 2, hp: 1, at: 1).

Units of type worker, light, heavy and ranged can move one cell in one of four directions. Units of player 1 are displayed with a blue outline, whereas units of player 2 are outlined in red. Movement is visualized with a gray line, attacking with a red line and gathering with a green line.




Figure 2. Screenshot basesWorkers16x16A map [Huft, 2025].

**Game Setup.** Most games are decided within 2000 steps. Therefore, the maximum number of steps is set to 2000. The experiments in this thesis were conducted using the RAISocketAI re-implementation of MicroRTS and the map basesWorkers16x16A, which is shown in figure 2. This symmetric map starts each player with one base, five resources, two workers and 50 resources located near the base [Goodfriend, 2024].

- **PassiveAI:** PassiveAI does not act.
- **WorkerRushAI and LightRushAI:** These bots build their corresponding units as quickly as possible and rush the opponent with them. WorkerRushAI is later used as a training opponent in the PPO fine-tuning phase after LT.
- **RandomAI and RandomBiasedAI:** RandomAI and RandomBiasedAI are stochastic bots. RandomBiasedAI is more structured than RandomAI and therefore a stronger opponent.
- **CoacAI:** CoacAI is a strong competition bot. It is used in the BC/PPO training of the Base Agent, later as a bot opponent for the Main Agent and also in the PPO fine-tuning phase. This bot uses ranged units, defended by heavy units. It tries to keep the opponent at a distance and to attack with ranged units, while the heavy units protect the ranged units from close combat.
- **Mayari:** Mayari is another strong scripted competition bot and is used together with CoacAI

in Base-Agent training, LT bot environments and fine-tuning. This bot primarily uses heavy units to rush the opponent.

- **Rojo:** rojo is a bot that tries to build many worker units and attack strategically with them.
- **MixedBot, Izanagi, Tiamat, Droplet, GuidedRojoA3N and NaiveMCTS:** These are additional stronger competition bots that are used for evaluation in order to test generalization for stronger bots. These bots are too slow to be used for training.

**Action Space.** The action space used in this thesis is implemented using the combination of GridNet and invalid action masking [Huang et al., 2021]. GridNet produces a unit action (an action for one specific cell) for each cell of the map, regardless of whether a unit is located on the cell or whether the action is valid [Han et al., 2019]. Building on this, invalid action masking is used, thereby also masking actions of empty cells [Huang et al., 2021]. The unit action space is defined as an 8-dimensional vector per cell:

$$(\text{pos}, \text{type}, d_{\text{move}}, d_{\text{harvest}}, d_{\text{return}}, d_{\text{produce}}, \text{pType}, \text{relPos})$$

The components of the unit-action vector are defined as follows:

- **pos:** linearized index of the cell in the grid (i.e., the grid represented in a one-dimensional space (e. g. for a  $16 \times 16$  grid: range:  $[0, 255]$ )),
- **type:** action type (range:  $[0, 5]$ , 0: NOOP, 1: Move, 2: Harvest, 3: Return, 4: Produce, 5: Attack),
- $d_{\text{move}}, d_{\text{harvest}}, d_{\text{return}}, d_{\text{produce}}$ : directions (range:  $[0, 3]$ , 0: North, 1: East, 2: South, 3: West) for the respective actions,
- **pType:** type of the unit to be produced (range:  $[0, 6]$ , 0: resource, 1: base, 2: barracks, 3: worker, 4: light, 5: heavy, 6: ranged),
- **relPos:** relative attack position (range:  $[0, 48]$ ).

For our agent, the unit action space (without the position dimension) is linearized and represented using one-hot encoding. The position is implicitly encoded by the linear grid indexing of the GridNet output [Huang et al., 2021]. Given a map of size  $h \times w$ , the action space has shape  $h \times w \times 78$ . For our  $16 \times 16$  map, the action space per environment in the GridNet representation would therefore contain  $16 \times 16 \times 78 = 19968$  actions. However, many of them can be invalid and are masked by invalid action masking [Huang et al., 2021]. Before applying the invalid action mask, the position is added to each action as a fixed



position index. Otherwise, masking would remove the grid structure of the action space, resulting in the removal of the implicit encoding for the position. Depending on the action type, only parts of the action vector are relevant (e. g. if the action type is `Move`, only the components `pos`, `type` and `dmove` are relevant). The other components are ignored.

**Observation Space.** Observations are represented as a grid of shape  $h \times w \times c$  per environment, where  $h$  and  $w$  are the height and width of the map and  $c$  is the number of encoded features of each cell in the game state (e. g. unit types, hit points, move direction, etc.) [Huang et al., 2021]. The screenshot of the map `basesWorkers16x16A` in figure 2 visualizes the observation space without any additional information.

### 3.2. Self-Play (SP)

In SP, the agent trains against its own current version, which leads to a continuous learning process with a co-evolving opponent [Silver et al., 2018]. This does not require additional diverse or strong opponents.

**Implementation.** To implement SP in the MicroRTS-Py re-implementation, Huang et al. implemented the environment to output the observations (obs) and rewards for both players, which allows the agent to train against itself by simply using the same policy for both players [Huang et al., 2021]. In RAISocketAI, that approach was improved and fixed [Goodfriend, 2024]. In this thesis, the same approach with some adjustments is used to implement SP, which is also used as a basis for FSP and LT.

#### Problems.

- The training is often unstable because of cyclical learning. For example, if policy B performs well against A, C performs well against B and A performs well against C, the agent may be trapped in a cycle in which it repeatedly relearns only the policy that performs well against itself, but not against other opponents [Vinyals et al., 2019].
- A big problem in the SP training setup is that the agent has to learn to play as player 2 as well as player 1. This complicates the training process and can lead to a slow learning process [Huang et al., 2021]. The biggest difference between player 1 and player 2 is that player 1 starts at the top left of the map, whereas player 2 starts at the bottom right. This means that the agent has to learn to start play in two different areas of the map, which can be difficult and time-consuming. The problem is further exacerbated by the fact that

the training of this thesis does not start from scratch. Rather it starts from a Base Agent that was only trained as player 1. Therefore, the agent would need to unlearn its bias towards starting at the top left of the map and learn to start at the bottom right as well.

**Solutions.** FSP/PFSP (3.4) are used to stabilize training and reduce cyclical learning. To address the second problem, the observations for player 2 are transformed by rotating them by 180 degrees. From the agent’s perspective, player 2 also starts at the top-left of the map. This makes the observations for player 1 and player 2 more similar and therefore makes SP with a pre-trained Base Agent possible, even though the Base Agent was exclusively trained as player 1 (figure 11). Actions are transformed accordingly, as a result the environment still executes the intended actions for player 2. Masks must be transformed as well, since they are based on the environment’s observations. The same rotation used for the observations is also applied to the invalid action masks. After masking, the selected actions are transformed back to fit the MicroRTS-Py environment. This includes a position component that is added before masking, since masking removes the implicit positional structure of the GridNet action space.

### 3.3. Fictitious Self-Play (FSP)

To address the instabilities of pure SP, in FSP the agent is trained not exclusively against the current policy, but against a uniformly sampled pool of the agent’s historical policies [Vinyals et al., 2019]. In the implementation of this thesis, this means that instead of player 2 always using the current policy, player 2 is sampled uniformly from a pool of Historical Agents (i. e. snapshots of the Main Agent at different training stages). This results in the learning-agent learning an effective response to the averaged opponent strategy, reducing cyclical learning. Under standard assumptions, FSP can converge towards an approximate Nash equilibrium in two-player zero-sum games (with respect to the average strategies). Heinrich et al. also demonstrate this empirically in poker experiments [Heinrich et al., 2015].

**Problem with the Implementation.** In the described implementation, player 1 has priority in conflict resolution (e. g. when both players want to move to the same cell). This gives player 2 a disadvantage, although it has only a minor effect on the results. The actual problem observed during training is that the agent can exploit this prioritization in FSP by optimizing the training objective primarily through improving player 1, while deliberately performing poorly as player 2. After sufficient training, the agent will learn a strategy that is based on player 1 being prioritized. Training then stagnates, since the agent already

has a high win rate against itself without developing a robust strategy against other opponents.

**Example.** The agent could learn to block an opponent in front of the base by issuing a simultaneous move to the same target cell. In FSP, player 2 would then be blocked, while player 1 would not, allowing the agent to achieve a win rate close to 100% against itself.

**Relevance for LT.** In LT, the result of the player 1 priority is not quite as problematic, because there are exploiters that can be trained against, which remain adversarial to the Main Agent. However, a large portion of the training becomes ineffective and provides poor learning signals if the agent exploits the player 1 priority. On top of that, this problem is particularly relevant for LT, which combines SP and FSP, because only Historical Agents are available for training in FSP, making the feedback for player 2 very sparse. In LT the Main Agent can directly train against the current Main Agent. This yields more immediate feedback on the Main Agent’s performance as an opponent (player 2). As a result, the agent no longer has to plan as far ahead in order to exploit this priority and minimize its own win rate as player 2.




Figure 3. Screenshot showing an action conflict in the game. Both heavy units of both players wanted to move to the same cell, as the opponent’s heavy units.

**Solutions.** In the re-implementation of MicroRTS-Py, there is a mode that resolves conflicts in a randomized manner. However, this does not help in this case, because

even before the conflict arises, invalid action masking masks out the action of player 2 since it is invalid. To solve this problem, the learning agents in the parallel games are alternated between playing as player 1 and player 2.

### 3.4. Prioritized Fictitious Self-Play (PFSP)

LT uses an extended variant of FSP called PFSP [Vinyals et al., 2019]. PFSP extends FSP by selecting opponents against which the agent can train efficiently, rather than selecting opponents uniformly. The selection probability is based on the win rate of the learning agent against the potential opponent [Vinyals et al., 2019]. As in AlphaStar, the PFSP win rate in this thesis is calculated by considering a draw to be 50% of a win for the learning agent. Win rates outside of PFSP are calculated without considering draws.

**Prioritization Function.** The exact prioritization function can differ for different league agent types. In AlphaStar, for the Main Exploiters PFSP prioritizes opponents of medium difficulty (PFSP win rate close to 50%). While the Main Agent prioritizes the strongest opponent (against which the PFSP win rate is closest to 0%). Adapted prioritization functions tailored to our training setup are used in this thesis, which deviate from those used in AlphaStar. An important difference is that training with an opponent against which the Main Agent has a win rate (without draws) close to 0% yields little informative gradient or learning signal.

**Implemented Prioritization Function.** To avoid prioritizing opponents that provide little learning signal, the following prioritization function was used for the Main Agent:




Figure 4. Prioritization function for the Main Agent in PFSP.

As can be seen in the figure 4, we prioritize opponents with a PFSP win rate around 50%. That is because most of our matchups are very draw-heavy. Therefore, it is likely that the win rate (without draws) is close to 0%, when the PFSP



win rate is below 50%. As a result, such opponents would otherwise be selected more frequently despite providing only a very weak learning signal. Furthermore, opponents with a PFSP win rate of 0% are prioritized over those with a PFSP win rate of 100%, since there is substantial learning potential for an opponent with a PFSP win rate of 0%, whereas a PFSP win rate of 100% usually means that the agent already defeats the opponent reliably. In this thesis PFSP was modified in a way that all agents have a non-zero probability of being selected. This ensures that no opponent is completely ignored because it lost all initial games (yielding a 100% PFSP win rate for the learning agent), since these initial results may have been outliers.

### 3.5. League Training (LT)

LT combines and extends SP and PFSP to further stabilize training. It also promotes robustness by training against opponents sampled from a league, including Main Exploiters and League Exploiters and Historical Agents [Vinyals et al., 2019].

**League Composition and Roles.** Learning agents of the league were trained in parallel. Similar to AlphaStar [Vinyals et al., 2019], the league in this thesis consists of the following league agent types:

- **Main Agent:** The primary policy that is optimized and evaluated.
- **Historical Agents:** Frozen snapshots of earlier Main Agents or exploiter agent policies. They stabilize training by providing a distribution of past strategies and mitigate catastrophic forgetting.
- **Main Exploiters:** Agents trained only against historical Main Agents, which encourages them to identify and exploit weaknesses in the Main Agent.
- **League Exploiters:** Agents trained against all Historical Agents, which further promotes diversity by identifying systemic weaknesses of the league.

Historical Agents are the only league agent type that is not a learning agent. The main differences between the roles are the opponent-sampling distribution and the hyperparameters. All agents share the same observation/action representation (GridNet with invalid action masking) and use the same underlying PPO algorithm.

**Training Loop and Snapshot Management.** LT iterates through cycles consisting of the following phases:

- **Select opponent:** sample an opponent from the league,

- **Collect rollouts:** run matches between the learning agent and the sampled opponent for a fixed number of steps, no agent is updated during this phase,
- **Update policy:** perform PPO updates using the collected trajectories against the sampled opponent,
- **Update statistics:** periodically evaluate against a subset of league members to update win rate estimates used by PFSP,
- **Add snapshots:** at predefined intervals (or when performance criteria are met), store a frozen copy of the current agent as a new Historical Agent.

An agent proceeds to the next phase only after every league agent type (Main Agent, Main Exploiters, League Exploiters) has completed the current phase.

**Bot Environments.** Bot environments are added to the league to stabilize training by continuing the PPO training against the bots CoacAI and Mayari, which the Base Agent was trained against. Only the Main Agent trains against the bots, since learning signals from bots could interfere with identifying and exploiting weaknesses in the Main Agent or the league, which is the main purpose of the exploiters (e. g. if the bots are too strong against an exploit that the Main Agent has not yet trained against). In MicroRTS-Py the environments for bots and for SP are not interchangeable. This is because the bots are not implemented as neural policies and therefore cannot be used in SP/FSP/LT directly. The number of parallel environment instances and their built-in bot opponent (if any) have to be defined at initialization of the MicroRTS-Py environments, which means the number of environments for each built-in bot remains fixed. One of the main advantages of LT is that training against opponents with bad learning signals is reduced. To retain that improvement we want the number of bot environments to be dynamic. Therefore, the number of bot environments is adjusted dynamically by reinitializing a MicroRTS instance, which is only used for training against the bots and not for SP/FSP/LT (because reinitialization also means resetting the environments in the MicroRTS instance). This MicroRTS instance for bots executes its step function at the same time as the MicroRTS instance for SP/FSP/LT.

**Number of Parallel Bot Environment Instances.** Two hyperparameters define the maximum number of parallel bot environment instances and the minimum number of parallel bot environment instances. Among other things, references to the observation, action and invalid action mask buffers must be adjusted to the correct size and partially reset whenever the number of parallel bot environment instances is changed.

**Summary: Opponent Selection.** PFSP samples the next opponent from a given set of possible opponents, with probabilities based on the current learning agent’s win rates against the potential opponents, whereas SP always trains against the current learning agent. The opponent selection for LT is illustrated in Figure 10. Depending on the agent role, opponents are either the current Main Agent or sampled from specific subsets of Historical Agents using PFSP. The Main Agent also trains against additional bot environments, while Main Exploiters and League Exploiters are matched only against agents from the league.

**Resetting Exploiters.** The reset procedure for exploiters during rollout collection is illustrated in figure 5




Figure 5. When resetting exploiters skip the update for the current rollout.

**AlphaStar LT Pseudocode.** In addition to the modifications made to adapt the LT pseudocode to the MicroRTS environment the following corrections to the AlphaStar pseudocode [Vinyals et al., 2019] were made:

- When resetting exploiters, the pseudocode first resets the weights and only then saves the snapshot for the Historical Agent. In our implementation, this order is reversed so that the snapshot is not taken from the reset exploiter.
- Main Exploiters were only trained against the current Main Agent, if they had a high win rate against the Main Agent from the very beginning. As soon as this win rate dropped, the Main Exploiter stopped training against the Main Agent for the remainder of this reset cycle. However, in the pseudocode Main Exploiters were only reset/checkpointed (creating a historical Main Exploiter) if they had a sufficiently high win rate against the Main Agent. This means that a Main Exploiter would never reset if it initially had a

low win rate against the Main Agent. Since the current Main Exploiter has already trained against the Base Agent, it is very likely that it will never checkpoint before the maximum number of steps is reached, even if it has a high win rate against all historical Main Agents. To solve this problem, the win rate used for resetting or checkpointing a Main Exploiter is defined as

$$w_{\text{reset}} = \max\left(\min_i w_{\text{hist},i}, w_{\text{main}}\right),$$

where  $w_{\text{hist},i}$  denotes the win rate against historical Main Agent  $i$  and  $w_{\text{main}}$  the win rate against the current Main Agent. This ensures that the Main Exploiter resets and checkpoints as soon as it achieves sufficiently high win rates against the available historical Main Agents (making high win rate against the Main Agent likely) and no longer has much to learn from the available opponents.

- Another change to the pseudocode is that win rates are set to 0.5 for the initial games. This makes it less likely that the agent will ignore opponents it has not yet played against because of outlier results with very high or very low win rates, including the Main Exploiter versus Main Agent matchup.

### 3.6. Training Pipeline.

The training pipeline extended in this thesis consists of four phases: BC, BC-FT, PPO and LT (or SP/FSP). The first three phases were already implemented in the previous work [Huft, 2025] and are only briefly described here.

- **BC, BC-FT:** The BC phases reduce the initial, often inefficient exploration in DRL methods such as PPO, which can be very time-consuming when trained from scratch [Torabi et al., 2018, Martin and Jhala, 2021]. The experts used for BC are CoacAI and Mayari [Huft, 2025].
- **PPO:** The third phase improves the generalization and performance of the agent through additional exploration beyond expert demonstrations [Martin and Jhala, 2021]. PPO constrains the size of policy updates in order to achieve stable training [Schulman et al., 2017].
- **LT:** Using LT, we aim to further generalize the Base Agent by training PPO not only against the MicroRTS bots (CoacAI and Mayari) (against which the Base Agent already achieves a high win rate), but also against opponents of equal strength or stronger. However, many opponents at this level of performance are not suitable for training, since they are very slow and thus slow down the collection of rollouts for PPO. In addition, the agent’s generalization capability

improves only to a limited extent when training against only a few new opponents, if the training environment does not provide sufficient opponent diversity. LT addresses both problems by providing numerous opponents that continue to evolve over the course of training and achieve a high win rate against historical versions of the current agent.

### 3.7. Additional Modifications for Diversity

The biggest problem with the training setup is the lack of opponent diversity. To make exploiters more diverse, exploration can be increased with hyperparameters such as the learning rate, the entropy coefficient and the Kullback-Leibler (KL) regularization coefficient.

**Annealed Entropy Coefficient.** The annealed entropy coefficient decreases the weight of the entropy bonus in the loss function and thereby decreases the exploration of the agent over the course of training. This was implemented for the exploiters to further increase exploration during the early training steps, while retaining a more deterministic policy in the later stages of training. This coefficient resets when the corresponding exploiter resets. In the case of a League Exploiter, it is possible that the exploiter checkpoints but does not reset. If that happens, the annealed entropy coefficient is reinitialized to half of its previous initial value.

**New Initial Agents.** To mitigate insufficient opponent diversity, we created new initial agents that focused on different unit types and can be used as alternative initial agents for the Main Exploiters in LT. We did not create new initial agents for the Main Agent because we prioritized further training the already strong Base Agent.

### 3.8. Delta Score.

To better guide the agent in the sparse reward landscape of MicroRTS, we added the change in the MicroRTS score between two consecutive time steps to the reward [Huang et al., 2021, Ontañón, 2013]. The score for each environment instance is calculated based on the current game state, which includes the number and type of units, their hit points and the resources of both players. The delta score is the difference between the score at the current time step and the score at the previous time step, which makes it a useful immediate learning signal for the agent. The score does not always accurately predict how well an agent is playing, which is why a coefficient is used to scale the delta score, so that the accumulated weighted delta score over a game is much smaller than the win/loss reward. In the prior work by M. Huft [Huft, 2025], the agent we now use as the Base Agent was trained with an attack bonus,

which is a reward for attacking the opponent’s units. The attack bonus was removed for our training and we adjusted the per-step delta score in the early stages of training to be approximately as large as the per-step attack bonus used in earlier training runs.

## 4. Experiments.

We use the Base Agent from the prior work of M. Huft [Huft, 2025], which was trained by using a combination of BC and PPO, as the initial agent in this thesis unless otherwise specified. The map used is `basesWorkers16x16A`, a symmetric map also used in related work such as [Huang et al., 2021, Huft, 2025, Goodfriend, 2024]. We evaluate a version of the Main Agent, which was not trained for the full amount of total steps. Our Main Agent is trained in LT with 4 Main Exploiters, 1 League Exploiter, with 25 environments each, for approximately 130000000 total steps and 1600 updates over 100 hours. The Main Agent played around 4500 games against the bots CoacAI and Mayari. During the training, win rates against both bots over the global training steps were monitored and shown in figure 6. In the training, the bot environments were evenly split between the two bots, with a maximum of 10 and a minimum of 5 environments for both bots combined. The number of bot environments throughout training is shown in figure 7. This results in a total of 160 parallel environments for LT when all bot environments are active. It requires a lot of computational resources, especially since the observations are not purely boolean and therefore required more memory. The peak memory usage increases by approximately 2500 MiB for an agent with 25 parallel environments. However, it is important to keep the number of environments per agent above a certain threshold, because it affects the rollout size and a too small rollout can lead to unstable training.




Figure 6. Win rates of CoacAI and Mayari against the Main Agent over the course of training.




Figure 7. Number of parallel bot environments over the course of training.

### 4.1. Hyperparameters

Notable hyperparameters for the Main Agent include a higher learning rate of  $5e-5$  than the Base Agent, namely  $2.5e-5$ , as well as higher value function coefficients of 0.15 for the Main Agent and 0.1 for the exploiters. An annealed entropy coefficient is used, starting at 0.025 and ending at 0.0075. The annealed entropy coefficient for exploiters is reset when the exploiter resets, as explained in 3.7. After sufficient steps, the entropy coefficient for exploiters remains high because of repeated resets, which helps them find exploits faster, while the annealing helps them retain and further improve these exploits. The entropy coefficient is very high in comparison to that of the Base Agent, which uses a constant coefficient of 0.005. This was mostly done to overcome the strong preferences of the Base Agent, which can make it more difficult for the agent to change its strategy and explore. All hyperparameters used are listed in Table 10.

### 4.2. Evaluation

We evaluated the Main Agent against the MicroRTS bots and the Base Agent, across at least 100 games per opponent. The opponents against which the evaluation is conducted include the MicroRTS built-in bots consisting of PassiveAI, WorkerRushAI, LightRushAI, RandomAI and RandomBiasedAI. Tables 2 and 3 below give a compact overview of the win rates and loss rates of the evaluated agents against all bot opponents. The detailed evaluation results of the Base Agent are listed below in Table 5 for comparison. The results vary from those reported in the thesis of M. Huft [Huft, 2025] because the original Base Agent from the prior work could not be found. Therefore a different version of the Base Agent provided by M. Huft, but trained using the same procedure, was used. The detailed evaluation results of the Main Agent are presented below in Table 6.

**Checkpoint and Reset Statistics.** In total there are 7 checkpoints for the Main Agent, 9 checkpoints for the Main Exploiters and 2 checkpoints for the League Exploiter. Table 4 presents the number of checkpoints for each initial agent type and the approximate total games that exploiters trained with them.

### 4.3. Results

The agent learned to keep the base defended, even while attacking with other units, which is a clear improvement over the Base Agent, which often leaves the base undefended while attacking. As shown in the evaluation results in Tables 2 and 3, LT improves the agent’s performance against almost all evaluated opponents. The detailed differences between the win rates are listed in Table 7 for better comparison. The win rates against the bots used in training reach 100%, which is a major improvement over the Base Agent, which achieves a 96% win rate against CoacAI, but only a 67% win rate against Mayari. The results of our Base Agent differ substantially from the results reported in the thesis of M. Huft [Huft, 2025]. However, the Main Agent still shows improvement over the results reported in that thesis, which report a 96% win rate against CoacAI and a 95% win rate against Mayari. The trained Main Agent also performs well against the Base Agent against which it was trained, achieving a win rate of 89%. The win rate against RandomBiasedAI is another major improvement, increasing from 79% to 97%. This indicates that, in comparison to the Base Agent, the trained Agent performs substantially better against bots that try to draw the game instead of win the game. The agents in the section “Other Bots” are considerably stronger against the Base Agent than the built-in bots, which is reflected in the lower win rates against them. However, the Main Agent still shows improvements over the Base Agent against most of these opponents. Highlights include the win rate against Rojo, which improves from 81% to 100%, while the win rate against MixedBot also increases from 43% to 57%. The win rate against Izanagi falls from 72% to 70%, however the loss rate improves from 16% to 7%, suggesting that the Main Agent defends the base more effectively while attacking. WorkerRushAI is the only opponent against which the Main Agent performs notably worse than the Base Agent. This shows that the Main Agent still has weaknesses that were not found.

Table 2. Comparison of win rates against all bot opponents. All values are given in %. Higher values are better.

| Opponent                     | Main Agent<br>without unit focus | Base Agent   | Main Agent (LT) | Main Agent +<br>Finetuning |
|------------------------------|----------------------------------|--------------|-----------------|----------------------------|
| <i>bots used in training</i> |                                  |              |                 |                            |
| coacAI                       | <b>100.0</b>                     | 96.0         | <b>100.0</b>    | <b>100.0</b>               |
| mayari                       | <b>100.0</b>                     | 67.3         | <b>100.0</b>    | <b>100.0</b>               |
| <i>Built-in bots</i>         |                                  |              |                 |                            |
| workerRushAI                 | 31.0                             | <b>100.0</b> | 35.0            | 99.5                       |
| passiveAI                    | <b>100.0</b>                     | 99.7         | <b>100.0</b>    | <b>100.0</b>               |
| lightRushAI                  | <b>100.0</b>                     | 98.0         | <b>100.0</b>    | <b>100.0</b>               |
| randomAI                     | <b>100.0</b>                     | 99.0         | <b>100.0</b>    | <b>100.0</b>               |
| randomBiasedAI               | 91.0                             | 79.0         | 97.0            | <b>97.5</b>                |
| <i>Other bots</i>            |                                  |              |                 |                            |
| rojo                         | <b>100.0</b>                     | 81.0         | <b>100.0</b>    | <b>100.0</b>               |
| mixedBot                     | 45.0                             | 43.3         | 57.0            | <b>60.5</b>                |
| izanagi                      | 68.0                             | 71.7         | 70.0            | <b>72.5</b>                |
| tiamat                       | 94.0                             | 93.0         | <b>99.0</b>     | <b>99.0</b>                |
| droplet                      | 45.0                             | 67.7         | 62.0            | <b>75.5</b>                |
| guidedRojoA3N                | 64.0                             | 74.7         | 87.0            | <b>88.5</b>                |
| naiveMCTS                    | 0.0                              | 0.7          | 2.0             | <b>6.5</b>                 |
| overall                      | 74.1                             | 76.5         | 79.2            | <b>85.7</b>                |

Table 3. Comparison of loss rates against all bot opponents. All values are given in %. Lower values are better.

| Opponent                     | Main Agent<br>without unit focus | Base Agent | Main Agent (LT) | Main Agent +<br>Finetuning |
|------------------------------|----------------------------------|------------|-----------------|----------------------------|
| <i>bots used in training</i> |                                  |            |                 |                            |
| coacAI                       | <b>0.0</b>                       | 2.0        | <b>0.0</b>      | <b>0.0</b>                 |
| mayari                       | <b>0.0</b>                       | 27.3       | <b>0.0</b>      | <b>0.0</b>                 |
| <i>Built-in bots</i>         |                                  |            |                 |                            |
| workerRushAI                 | 63.0                             | <b>0.0</b> | 55.0            | 0.5                        |
| passiveAI                    | <b>0.0</b>                       | <b>0.0</b> | <b>0.0</b>      | <b>0.0</b>                 |
| lightRushAI                  | <b>0.0</b>                       | 2.0        | <b>0.0</b>      | <b>0.0</b>                 |
| randomAI                     | <b>0.0</b>                       | <b>0.0</b> | <b>0.0</b>      | <b>0.0</b>                 |
| randomBiasedAI               | <b>0.0</b>                       | <b>0.0</b> | <b>0.0</b>      | <b>0.0</b>                 |
| <i>Other bots</i>            |                                  |            |                 |                            |
| rojo                         | <b>0.0</b>                       | 5.3        | <b>0.0</b>      | <b>0.0</b>                 |
| mixedBot                     | 33.0                             | 34.3       | 20.0            | <b>13.5</b>                |
| izanagi                      | 15.0                             | 16.3       | 7.0             | <b>6.0</b>                 |
| tiamat                       | 6.0                              | 7.0        | 1.0             | <b>0.5</b>                 |
| droplet                      | 22.0                             | <b>9.0</b> | 22.0            | 10.5                       |
| guidedRojoA3N                | <b>0.0</b>                       | 1.7        | <b>0.0</b>      | <b>0.0</b>                 |
| naiveMCTS                    | <b>0.0</b>                       | 2.7        | <b>0.0</b>      | <b>0.0</b>                 |
| overall                      | 9.9                              | 7.7        | 7.5             | <b>2.2</b>                 |

**Main Agent without Unit Focus Evaluation.** The detailed evaluation results in Table 8 show an example evaluation for a Main Agent trained without unit focus. The compact Tables 2 and 3 also include most of these results for better comparison. As can be seen in the table, the Main Agent with unit focus performs better or similar in all

matchups. The biggest differences are the win rates against Droplet, which improve with the unit focus from 45% to 62% and GuidedRojoA3N, which improves from 64% to 87%.



**PPO Fine-Tuning.** After the LT training, the Main Agent was further trained with PPO for approximately 9 million steps. The 25 environments were evenly split between CoacAI, Mayari and WorkerRushAI. The learning rate was set to  $1.5e-6$ , the entropy coefficient was set to 0.005 and no exploiters were used during this fine-tuning phase. The detailed results over 200 evaluation games per opponent are presented below in Table 9. The compact Tables 2 and 3 also show most of these results. After fine-tuning, the win rate against WorkerRushAI is now almost 100% and most other win rates also improved. The averaged win rate against all evaluation opponents excluding the Base Agent is 77% for the Base Agent. After LT and PPO fine-tuning, this averaged win rate increases to 86%, which is a substantial improvement of 9%. This averaged win rate for the reported results of the Base Agent in the thesis of M. Huft [Huft, 2025] is 79%, which means that the Main Agent after LT and PPO fine-tuning still shows an improvement of 7% over the results reported in that thesis. If we also exclude the bots WorkerRushAI and PassiveAI, for which the Base Agent already has a win rate of approximately 100%, the averaged win rate against the remaining opponents increases from 73% to 83%, which is an improvement of almost 10%. The total loss rate against all opponents excluding the Base Agent is decreased from 8% to 3%.

#### 4.4. Discussion

LT improved the Main Agent’s win, draw and loss rates against a wide range of opponents, including the MicroRTS bots and the Base Agent used during training and generalized to other bots that were not used in training. Our results suggest that LT can improve the Main Agent’s performance further than the BC/PPO training was able to, without obvious signs of overfitting to the opponents used in training even after many different training sessions and further PPO training [Huft, 2025]. Our interpretation is that pure PPO does not search for well-generalizing policies. Instead, the training optimizes performance against the limited set of training opponents. It therefore may pass only briefly through policies that generalize well, before continuing towards more specialized solutions. In contrast, LT changes the training opponent distribution such that robust performance against diverse opponents becomes part of the optimization target itself. This makes it more likely that training explores and retains a broader range of policies with good generalization. The results show that LT was able to discover policies that generalize well while also outperforming the Base Agent even against its original training opponents despite being trained against a larger set of opponents.

**PPO Fine-Tuning.** After LT, additional PPO fine-tuning further improved the Main Agent’s performance, especially when training against WorkerRushAI. WorkerRushAI probably exploited a weakness of the Main Agent that was not found during LT. That is presumably the reason for the win rate against the other opponents improving after fine-tuning, since the agent learned to defend against this exploit. These improvements also show that the fine-tuning did not overfit to the training opponents and instead generalized to other opponents. The lower loss rate suggests that the agent improved not only in attacking the opponent, but also at defending the base.

**Player 1 Priority.** Alternating the learning agents between playing as player 1 and player 2 does not prevent the prioritization of player 1. However, evaluating an agent against itself suggests that the prioritization of player 1 may have no major impact on the results. For example, one evaluation resulted in a win rate of 38% for player 1, 28% draws and 34% wins for player 2 over 200 evaluation games, which suggests that player 1 did not have a substantial advantage. At the same time, alternating the player side prevents the exploitation of the prioritization by the learning agent, since it can now also be player 2 and therefore a strategy that is only effective for player 1 is no longer beneficial.

**Diversity Definition.** A diverse agent is an agent with a broad effective strategy repertoire. This means that such an agent is capable of using multiple qualitatively different strategies against the same opponent while still performing well, rather than mainly relying on the same strategy against one opponent. Therefore, a diverse agent in this thesis is not merely one with variation in its behavior or one that reacts differently to different opponents. Instead, it is one that has the ability to represent and use substantially different strategic approaches against the same opponent.

In this thesis a set of agents can be diverse even if the individual agents in the set are not diverse themselves. A set of agents is considered diverse if the definition of a diverse agent would apply to the set as a whole.

**Base Agent Strategy.** Our Base Agent prefers a strategy that focuses on producing heavy units as fast as possible and rushing the opponent with them. We will call this a heavy rush. When pressured by worker units, the agent defends with its own worker units until it has produced the first heavy unit. In case of a lot of worker units, it also produces ranged units to attack the worker units from a distance.

**Base Agent Diversity.** The Base Agent strongly favors a single strategy focused on heavy units, as described in 4.4.

Much of the Base Agent’s strategic variation is reactive, so its play against similar variants shows even less strategic variation. The Base Agent is able to perform well against a broad set of opponents and can use different unit types if needed. However, the Base Agent itself is not diverse in the sense defined above in 4.4.

**Exploiter Diversity as the Main Limitation.** The exploiters are therefore less likely to cover those weaknesses that are exploitable from strategies different from their initial agents’ policies. This can be seen in the high win rates against e. g. Tiamat, which is a bot using heavy units. The agents are mostly trained against other agents that are initialized from the same Base Agent. We presume the limited diversity of the Base Agent to be the reason why the exploiters and the Main Agent tend to encounter similar opponents and therefore exhibit mostly similar games and strategies against each other. This suggests a main limitation of the training approach is that, because of the limited diversity of the Base Agent (4.4), the diversity of the set of historical Main Exploiters is also limited unless very long training times and a high degree of exploration are used. Some weaknesses of the Main Agent that can be exploited by strategies for different unit types are therefore very difficult to represent with historical Main Exploiters. Such strategies are difficult to discover through LT alone, because they are far from the Base Agent’s default heavy-rush behavior and often require first unlearning fundamental parts of the initial strategy before a substantially different strategy can be learned. In between the switching over to a unit type, the agent performs very poorly, which suggests that there are distant clusters of local maxima in the reward landscape corresponding to different unit types. In a complex environment such as MicroRTS, there are likely many more weaknesses that can be exploited by only slightly different strategies than there are historical exploiters that can be created during training.


**Effects of New Initial Agents.** For this reason, the introduction of new initial agents was particularly important for better win rates against diverse opponents. We hypothesize that the diversity of the initial agents is crucial for the diversity of the exploiters and possibly more important than the initial agents’ win rate. This is why the new initial agents, with a focus on different unit types, were implemented, allowing the exploiters to find exploits related to different unit types without first having to unlearn the strategy for the unit type favored by the Base Agent. The new initial agents contributed to generalization across diverse opponents. Tables 2 and 3 show that performance improved against a wider range of opponents compared to the previous evaluation that did not use the new initial agents. One example is the win rate against Droplet,

which improves from 45% to 62%, while the win rate against GuidedRojoA3N improves from 64% to 87%.

The exploiter win rates and unit compositions against the Main Agent suggest that exploiters are capable of identifying and exploiting weaknesses within the set of strategies for the respective unit type favored by the Base Agent. This also suggests that pure LT is likely to work efficiently only if the Base Agent is already diverse in the sense defined above. At the same time, figure 9 for the worker-focused initial agent shows that at least some mixed strategies can be achieved with this training setup. This suggests that with the new initial agents focused on different unit types, the historical exploiter set is diverse enough to cover many relevant strategy variations and improve the performance of the Main Agent, even against many other agents that were not included in the training or evaluation. However, the result against WorkerRushAI suggests that more diverse exploiters could further improve the Main Agent’s performance.

Figure 8 summarizes this relation schematically.

### Without new initial agents



### With new initial agents



Figure 8. Illustration of our hypothesis about how the diversity of the Base Agent and initial agents affects exploiter diversity and the Main Agent’s generalization during LT.

**Comparison to AlphaStar.** This diversity limitation was probably less problematic in AlphaStar [Vinyals et al., 2019], because its initial policy was pre-trained by imitation learning on human replay data, which covered a much broader range of strategies. Therefore, the initial policy was likely competent across a broader range of strategies than the Base Agent used in this thesis, which was trained with BC and PPO against only two bots and was therefore biased towards a comparatively narrow strategy set. This broader initialization presumably made it easier for exploiters to reach and exploit a more diverse set of strategies. The diversity of our Base Agent is further constrained in comparison to AlphaStar by the fact that, during BC and PPO, some possible opponents must be excluded from

training in order to obtain a more robust evaluation, while many strong agents are too slow to be practical training opponents.

**Bot Environments.** The additional bot environments for the Main Agent are particularly important in the initial phase of LT, in which the Main Exploiters have not yet found weaknesses in the Main Agent. In this LT phase, the Main Agent can only train against the historical Main Agents and itself. The bots provide additional opponents that are not based on the Main Agent’s current strategy, hence reducing the risk of overfitting. After this initial LT phase, the bot environments are less important. That is why the number of bot environments is dynamic over the course of training and often relatively low as shown in figure 7.

**Other Limitations.** Despite the improved diversity and general performance of the Main Agent, there are still some limitations of the training approach:

- **Exploiter diversity and computational cost.** LT also requires a lot of computational resources. Our setup has a maximum of 160 parallel environments, while pure PPO only uses 24 parallel environments [Huft, 2025]. This causes a much higher memory usage and long training times.
- **Map constraints.** This thesis, just like the thesis for the Base Agent [Huft, 2025], only uses one map (basesWorkers16x16A), with full observability of all units and resources. This is a very specific setting.

## 4.5. Future Work.

Future work could focus on multiple different maps and on fog of war. Decreasing the computational cost of training is also an important direction for future work. Another possibility is to improve the diversity of the exploiters, which is a crucial factor for the performance of the Main Agent. This can be done by increasing diversity in either the initial agents, the Base Agents or the training itself.

**Map Diversity.** Map diversity would remove the map constraints and increase the diversity at the same time. Changing maps is already possible in MicroRTS-Py. Automatically generating new maps with different layouts and resource distributions and roughly equal chances of winning for both players is an interesting direction for future work, which could encourage the agent to learn more generalizable strategies and to adapt to different situations. This would also help to mitigate the problem of insufficient diversity if the maps force the agent to learn different strategies, for example by starting the training on maps that already contain units of different unit types, which

can encourage different strategies. Another promising idea would be to train an agent to generate symmetric maps on which the learning agent is weaker than the bot opponent during PPO or LT. This would automatically generate maps that are tailored to the learning agent’s weaknesses and could force the learning agent to improve in various strategies for different scenarios. The learning agent trained on those maps could then be used as a diverse initial agent for the exploiters and the Main Agent. Such a map generation approach could also be beneficial in LT.

**Initial Agents.** The current initial agents for the exploiters are trained mainly with BC with different experts, because of time constraints. Improving these initial agents to a similar level as the Base Agent, while focusing on generalization and diversity, could further improve the effectiveness of the LT method. Alternatively, a new strong Base Agent could be trained from scratch with a focus on diverse strategies, without using for example the PPO phase as implemented in this thesis, which resulted in a strong agent, that mostly favors one strategy for one unit type.

**MicroRTS efficiency.** A major bottleneck for training is the efficiency of the MicroRTS environment, which is currently implemented in Java and runs on the CPU. This is especially problematic for the large number of parallel environments used in LT. Reimplementing the environment on the GPU could substantially reduce training and evaluation times, which would make a more responsive training process and longer training runs within the same computational budget possible.

**Production Exploration Bonus.** To further increase exploration without increasing training time, an exploration bonus targeting the production of workers, ranged, heavy and light units can be added to the reward. This was achieved by using a mask that gives barracks and bases (units whose only function is to produce other units) an entropy bonus, which encourages more diverse production decisions and thereby increases exploration across different unit types, by making the agent interact with them. The production exploration bonus increased the variety of produced unit types, but it was not sufficient to make the agent explore strategies centered on different unit types.

**Unit Bonus.** This is a clipped unit bonus that can be added to the reward, encouraging the agent to use different unit types (e. g. by giving the Base Agent a bonus every time it produces a ranged unit). The number of other unit types increased when using the unit bonus. The unit bonus likewise increased the number of off-focus unit types, but even with weaker agents such as the agent after the BC

phase and over 20 million steps, the agent did not develop strategies centered on different unit types. Instead it often produced the bonus unit type only at the end of a round, when the result was already decided.

**Unit Bonus Distribution.** Unit bonus distribution is an improved method for increasing the diversity of exploiters based on the unit bonus. The idea is to continue training the agent with an additional input feature for the agent that specifies the unit bonus sampled from a distribution that is updated every game.

The unit bonus distribution allows the trained agent to be used as an exploiter with a static bonus or a distribution-based bonus, resulting in an exploiter with definable tendencies toward different unit types. This also allows the retraining of the PPO phase with the unit bonus, because only one agent needs to be trained instead of multiple agents with different fixed unit bonuses, which would be very time-consuming. Retraining the agent from scratch with the unit bonus distribution showed promising results. However, the training was not continued long enough to evaluate the effect of the unit bonus distribution on the diversity of the exploiters, since training with the unit bonus distribution is very time-consuming because the BC phase did not work with the current implementation and therefore had to be skipped. Starting the LT with the Base Agent and the unit bonus distribution instead showed the same problem as training with the unit bonus.

## 5. Conclusion.

In this thesis, we implemented a modified LT method from AlphaStar [Vinyals et al., 2019] in combination with PPO in the MicroRTS environment, using a Base Agent trained on BC and PPO as the initial agent for the Main Agent and the exploiters.

As shown in the results, LT improves the Main Agent’s performance against a wide range of opponents, including most of the MicroRTS bots and the Base Agent. Furthermore, LT allows the Main Agent to generalize to other bots that were not used in training, while also improving the win rates against the bots used in the original training of the Base Agent. This suggests that the combination of LT and PPO is an effective approach for improving the performance of a Base Agent trained on BC and PPO in MicroRTS.

Finetuning the Main Agent with PPO after LT further improved the Main Agent’s performance.

However, the diversity of the exploiters’ initial agents seems to be a crucial factor, which we were able to partially

address by introducing new initial agents with different unit type focuses. This solution has its own limitations, such as not guaranteeing that all relevant clusters in the reward landscape are covered by the initial agents and complicating the training process by requiring the training of multiple initial agents. In our approach, diverse bot agents for BC and PPO training therefore remain very important for the performance of the Main Agent. While exploiters can provide some diversity and reduce the need for many bot agents with similar strategies, they cannot fully replace diverse bot agents in this training approach. This is because substantially different strategies from the Base Agent are more difficult for exploiters to discover and therefore may not be represented.

This also suggests that starting LT from a BC-pre-trained agent may be suboptimal, especially when BC is based on only a small set of experts, because the BC phase can bias the initial agent toward a narrow set of strategies. As a result, it may become more difficult for the exploiters to discover strategically diverse responses, which can in turn limit the Main Agent’s robustness and generalization. Replacing our current BC and PPO approach with an approach that focuses on learning diverse strategies from the start could therefore further increase the Main Agent’s performance and generalization, while also reducing the need for many different bot agents for training.

In summary, this thesis demonstrates the effectiveness of LT for improving the performance of an agent trained on BC and PPO in MicroRTS, while also highlighting the importance of the diversity of the exploiters’ initial agents.






Figure 9. Number of units after each game for the agents initialized with worker-focused (shown are 2 full reset cycles).

Table 4. Checkpoint amount per initial agent type and the approximate total games that exploiters trained with them.

| Initial agent type | Main Agent | League Exploiter | Main Exploiter | total exploiter games |
|--------------------|------------|------------------|----------------|-----------------------|
| Base Agent         | 7          | 2                | 4              | 50000                 |
| worker_focused     | 0          | 0                | 2              | 11000                 |
| ranged_focused     | 0          | 0                | 1              | 387                   |
| light_focused      | 0          | 0                | 2              | 14000                 |
| Total              | 7          | 2                | 9              | 75387                 |

Table 5. Base Agent evaluation

| Opponent                     | Games | Wins | Draws | Losses | Win rate $\uparrow$ | Draw rate | Loss rate $\downarrow$ |
|------------------------------|-------|------|-------|--------|---------------------|-----------|------------------------|
| <i>bots used in training</i> |       |      |       |        |                     |           |                        |
| coacAI                       | 300   | 288  | 6     | 6      | 0.96                | 0.02      | 0.02                   |
| mayari                       | 300   | 202  | 16    | 82     | 0.673               | 0.053     | 0.273                  |
| <i>Built-in bots</i>         |       |      |       |        |                     |           |                        |
| workerRushAI                 | 300   | 300  | 0     | 0      | 1                   | 0         | 0                      |
| passiveAI                    | 300   | 299  | 1     | 0      | 0.997               | 0.003     | 0                      |
| lightRushAI                  | 300   | 294  | 0     | 6      | 0.98                | 0         | 0.02                   |
| randomAI                     | 300   | 297  | 3     | 0      | 0.99                | 0.01      | 0                      |
| randomBiasedAI               | 300   | 237  | 63    | 0      | 0.79                | 0.21      | 0                      |
| <i>Other bots</i>            |       |      |       |        |                     |           |                        |
| rojo                         | 300   | 243  | 41    | 16     | 0.81                | 0.137     | 0.053                  |
| mixedBot                     | 300   | 130  | 67    | 103    | 0.433               | 0.223     | 0.343                  |
| izanagi                      | 300   | 215  | 36    | 49     | 0.717               | 0.12      | 0.163                  |
| tiamat                       | 300   | 279  | 0     | 21     | 0.93                | 0         | 0.07                   |
| droplet                      | 300   | 203  | 70    | 27     | 0.677               | 0.233     | 0.09                   |
| guidedRojoA3N                | 300   | 224  | 71    | 5      | 0.747               | 0.237     | 0.017                  |
| naiveMCTS                    | 300   | 2    | 290   | 8      | 0.007               | 0.967     | 0.027                  |
| overall                      | 4200  | 3213 | 664   | 323    | 0.765               | 0.158     | 0.077                  |

Table 6. Main Agent evaluation

| Opponent                | Games | Wins | Draws | Losses | Win rate $\uparrow$ | Draw rate | Loss rate $\downarrow$ |
|-------------------------|-------|------|-------|--------|---------------------|-----------|------------------------|
| <i>used in training</i> |       |      |       |        |                     |           |                        |
| Base Agent              | 100   | 89   | 9     | 2      | 0.89                | 0.09      | 0.02                   |
| coacAI                  | 100   | 100  | 0     | 0      | 1                   | 0         | 0                      |
| mayari                  | 100   | 100  | 0     | 0      | 1                   | 0         | 0                      |
| <i>Built-in bots</i>    |       |      |       |        |                     |           |                        |
| workerRushAI            | 100   | 35   | 10    | 55     | 0.35                | 0.1       | 0.55                   |
| passiveAI               | 100   | 100  | 0     | 0      | 1                   | 0         | 0                      |
| lightRushAI             | 100   | 100  | 0     | 0      | 1                   | 0         | 0                      |
| randomAI                | 100   | 100  | 0     | 0      | 1                   | 0         | 0                      |
| randomBiasedAI          | 100   | 97   | 3     | 0      | 0.97                | 0.03      | 0                      |
| <i>Other bots</i>       |       |      |       |        |                     |           |                        |
| rojo                    | 100   | 100  | 0     | 0      | 1                   | 0         | 0                      |
| mixedBot                | 100   | 57   | 23    | 20     | 0.57                | 0.23      | 0.2                    |
| izanagi                 | 100   | 70   | 23    | 7      | 0.7                 | 0.23      | 0.07                   |
| tiamat                  | 100   | 99   | 0     | 1      | 0.99                | 0         | 0.01                   |
| droplet                 | 100   | 62   | 16    | 22     | 0.62                | 0.16      | 0.22                   |
| guidedRojoA3N           | 100   | 87   | 13    | 0      | 0.87                | 0.13      | 0                      |
| naiveMCTS               | 100   | 2    | 98    | 0      | 0.02                | 0.98      | 0                      |
| overall                 | 1500  | 1198 | 195   | 107    | 0.799               | 0.130     | 0.071                  |

Table 7. Difference in evaluation results between the Main Agent and the Base Agent (e. g. 0.04 means the Main Agent increases the win rate by 4% in comparison to the Base Agent).

| Opponent                     | Win rate ↑ | Draw rate | Loss rate ↓ |
|------------------------------|------------|-----------|-------------|
| <i>bots used in training</i> |            |           |             |
| coacAI                       | 0.04       | -0.02     | -0.02       |
| mayari                       | 0.327      | -0.053    | -0.273      |
| <i>Built-in bots</i>         |            |           |             |
| workerRushAI                 | -0.65      | 0.1       | 0.55        |
| passiveAI                    | 0.003      | -0.003    | 0           |
| lightRushAI                  | 0.02       | 0         | -0.02       |
| randomAI                     | 0.01       | -0.01     | 0           |
| randomBiasedAI               | 0.18       | -0.18     | 0           |
| <i>Other bots</i>            |            |           |             |
| rojo                         | 0.19       | -0.137    | -0.053      |
| mixedBot                     | 0.137      | 0.007     | -0.143      |
| izanagi                      | -0.017     | 0.11      | -0.093      |
| tiamat                       | 0.06       | 0         | -0.06       |
| droplet                      | -0.057     | -0.073    | 0.13        |
| guidedRojoA3N                | 0.123      | -0.107    | -0.017      |
| naiveMCTS                    | 0.013      | 0.013     | -0.027      |

Table 8. Example for Main Agent without unit focus evaluation

| Opponent                | Games | Wins | Draws | Losses | Win rate ↑ | Draw rate | Loss rate ↓ |
|-------------------------|-------|------|-------|--------|------------|-----------|-------------|
| <i>used in training</i> |       |      |       |        |            |           |             |
| Base Agent              | 100   | 85   | 15    | 0      | 0.85       | 0.15      | 0           |
| coacAI                  | 100   | 100  | 0     | 0      | 1          | 0         | 0           |
| mayari                  | 100   | 100  | 0     | 0      | 1          | 0         | 0           |
| <i>Built-in bots</i>    |       |      |       |        |            |           |             |
| workerRushAI            | 100   | 31   | 6     | 63     | 0.31       | 0.06      | 0.63        |
| passiveAI               | 100   | 100  | 0     | 0      | 1          | 0         | 0           |
| lightRushAI             | 100   | 100  | 0     | 0      | 1          | 0         | 0           |
| randomAI                | 100   | 100  | 0     | 0      | 1          | 0         | 0           |
| randomBiasedAI          | 100   | 91   | 9     | 0      | 0.91       | 0.09      | 0           |
| <i>Other bots</i>       |       |      |       |        |            |           |             |
| rojo                    | 100   | 100  | 0     | 0      | 1          | 0         | 0           |
| mixedBot                | 100   | 45   | 22    | 33     | 0.45       | 0.22      | 0.33        |
| izanagi                 | 100   | 68   | 17    | 15     | 0.68       | 0.17      | 0.15        |
| tiamat                  | 100   | 94   | 0     | 6      | 0.94       | 0         | 0.06        |
| droplet                 | 100   | 45   | 33    | 22     | 0.45       | 0.33      | 0.22        |
| guidedRojoA3N           | 100   | 64   | 36    | 0      | 0.64       | 0.36      | 0           |
| naiveMCTS               | 100   | 0    | 100   | 0      | 0          | 1         | 0           |
| overall                 | 1500  | 1123 | 238   | 139    | 0.749      | 0.159     | 0.093       |

Table 9. Evaluation after PPO fine-tuning

| Opponent                   | Games | Wins | Draws | Losses | Win rate $\uparrow$ | Draw rate | Loss rate $\downarrow$ |
|----------------------------|-------|------|-------|--------|---------------------|-----------|------------------------|
| <i>used in training</i>    |       |      |       |        |                     |           |                        |
| Base Agent                 | 200   | 185  | 15    | 0      | 0.925               | 0.075     | 0                      |
| coacAI                     | 200   | 200  | 0     | 0      | 1                   | 0         | 0                      |
| mayari                     | 200   | 200  | 0     | 0      | 1                   | 0         | 0                      |
| workerRushAI               | 200   | 199  | 0     | 1      | 0.995               | 0         | 0.005                  |
| <i>Other built-in bots</i> |       |      |       |        |                     |           |                        |
| passiveAI                  | 200   | 200  | 0     | 0      | 1                   | 0         | 0                      |
| lightRushAI                | 200   | 200  | 0     | 0      | 1                   | 0         | 0                      |
| randomAI                   | 200   | 200  | 0     | 0      | 1                   | 0         | 0                      |
| randomBiasedAI             | 200   | 195  | 5     | 0      | 0.975               | 0.025     | 0                      |
| <i>Other bots</i>          |       |      |       |        |                     |           |                        |
| rojo                       | 200   | 200  | 0     | 0      | 1                   | 0         | 0                      |
| mixedBot                   | 200   | 121  | 52    | 27     | 0.605               | 0.26      | 0.135                  |
| izanagi                    | 200   | 145  | 43    | 12     | 0.725               | 0.215     | 0.06                   |
| tiamat                     | 200   | 198  | 1     | 1      | 0.99                | 0.005     | 0.005                  |
| droplet                    | 200   | 151  | 28    | 21     | 0.755               | 0.14      | 0.105                  |
| guidedRojoA3N              | 200   | 177  | 23    | 0      | 0.885               | 0.115     | 0                      |
| naiveMCTS                  | 200   | 13   | 187   | 0      | 0.065               | 0.935     | 0                      |
| overall                    | 3000  | 2584 | 354   | 62     | 0.861333333         | 0.118     | 0.020666667            |

Table 10. Hyperparameters and control parameters used for the LT setup

| Category                                                                                 | Parameter                                | Value             | Description                                                                                           |
|------------------------------------------------------------------------------------------|------------------------------------------|-------------------|-------------------------------------------------------------------------------------------------------|
| <b>LT Setup</b>                                                                          |                                          |                   |                                                                                                       |
| League                                                                                   | num_main_exploiters                      | 4                 | Number of Main Exploiters                                                                             |
| League                                                                                   | num_league_exploiters                    | 1                 | Number of League Exploiters                                                                           |
| League                                                                                   | selfplay_ready_save_interval             | $5 \cdot 10^5$    | Step interval for potential league checkpoints (checkpoint if the win rate is sufficiently high)      |
| League                                                                                   | main_selfplay_save_interval              | $6 \cdot 10^6$    | Step interval for forced Main Agent checkpoints                                                       |
| League                                                                                   | main_exploiter_selfplay_save_interval    | $9 \cdot 10^6$    | Step interval for forced Main Exploiter checkpoints                                                   |
| League                                                                                   | league_exploiter_selfplay_save_interval  | $9 \cdot 10^6$    | Step interval for forced League Exploiter checkpoints                                                 |
| League                                                                                   | main_winrate_threshold                   | 0.7               | Win rate threshold for Main Agent checkpoints                                                         |
| League                                                                                   | main_exploiter_winrate_threshold         | 0.7               | Win rate threshold for Main Exploiter checkpoints                                                     |
| League                                                                                   | league_exploiter_winrate_threshold       | 0.7               | Win rate threshold for League Exploiter checkpoints                                                   |
| League                                                                                   | main_exploiter_vs_main_winrate_threshold | 0.3               | Threshold for Main Exploiters against the Main Agent                                                  |
| PFSP                                                                                     | pfsp_min_prob_factor                     | 0.25              | Minimum probability factor added to PFSP                                                              |
| PFSP                                                                                     | main_PFSP_prob                           | 0.5               | Probability of using PFSP with historicals for the Main Agent                                         |
| PFSP                                                                                     | main_SP_prob                             | 0.35              | Probability of SP for the Main Agent                                                                  |
| <b>Environments and Training Scope</b>                                                   |                                          |                   |                                                                                                       |
| Environments                                                                             | num_bot_envs                             | 5                 | Initial number of bot environments.                                                                   |
| Environments                                                                             | min_num_bot_envs                         | 5                 | Lower bound for the number of bot environments                                                        |
| Environments                                                                             | max_num_bot_envs                         | 10                | Upper bound for the number of bot environments                                                        |
| Environments                                                                             | min_bot_winrate                          | 0.955             | Lower win rate threshold for bot environment adjustment (if crossed num_bot_envs == max_num_bot_envs) |
| Environments                                                                             | max_bot_winrate                          | 0.98              | Upper win rate threshold for bot environment adjustment (if crossed num_bot_envs == min_num_bot_envs) |
| Environments                                                                             | num_main_envs                            | 25                | Number of parallel environments for the Main Agent                                                    |
| Environments                                                                             | num_envs_per_main_exploiters             | 25                | Number of environments per Main Exploiter                                                             |
| Environments                                                                             | num_envs_per_league_exploiters           | 25                | Number of environments per League Exploiter                                                           |
| Training                                                                                 | total_timesteps                          | 200 000 000       | Total number of training steps (the training evaluated stopped early at about 130000000 steps)        |
| Training                                                                                 | checkpoint_end_buffer_steps              | $3 \cdot 10^6$    | Exploiters cannot checkpoint<br>checkpoint_end_buffer_steps steps before reaching total_timesteps     |
| <b>Reward Function</b>                                                                   |                                          |                   |                                                                                                       |
| Reward                                                                                   | attack_reward_weight                     | 0.0               | Weight of the attack reward                                                                           |
| Reward                                                                                   | rewardscore                              | 0.002             | Scaling factor of delta score.                                                                        |
| <b>PPO Hyperparameters (Main Agent)</b>                                                  |                                          |                   |                                                                                                       |
| PPO Main Agent                                                                           | PPO_learning_rate                        | $5 \cdot 10^{-5}$ | Learning rate of the Main Agent                                                                       |
| PPO Main Agent                                                                           | num_steps                                | 512               | Number of steps per rollout                                                                           |
| PPO Main Agent                                                                           | gamma                                    | 1.0               | Discount factor                                                                                       |
| PPO Main Agent                                                                           | gae_lambda                               | 0.95              | Generalized Advantage Estimation (GAE) parameter                                                      |
| PPO Main Agent                                                                           | main_ent_coef_min                        | 0.0075            | minimum entropy regularization coefficient                                                            |
| PPO Main Agent                                                                           | main_ent_coef_max                        | 0.025             | maximum entropy regularization coefficient                                                            |
| PPO Main Agent                                                                           | vf_coef                                  | 0.15              | Weight of the value loss                                                                              |
| PPO Main Agent                                                                           | max_grad_norm                            | 0.5               | Gradient clipping threshold                                                                           |
| PPO Main Agent                                                                           | clip_coef                                | 0.15              | PPO clipping coefficient                                                                              |
| PPO Main Agent                                                                           | update_epochs                            | 4                 | Number of update epochs per PPO cycle.                                                                |
| PPO Main Agent                                                                           | kle_stop                                 | true              | Early stopping when KL divergence becomes too large                                                   |
| PPO Main Agent                                                                           | kle_rollback                             | false             | No rollback when the KL threshold is exceeded                                                         |
| PPO Main Agent                                                                           | target_kl                                | 0.01              | Target KL divergence value                                                                            |
| PPO Main Agent                                                                           | kl_coef                                  | 0.3               | Coefficient of the KL penalty term                                                                    |
| PPO Main Agent                                                                           | norm_adv                                 | true              | Advantages are normalized                                                                             |
| PPO Main Agent                                                                           | anneal_lr                                | true              | Learning rate annealing is enabled                                                                    |
| PPO Main Agent                                                                           | clip_vloss                               | true              | Value loss clipping is enabled                                                                        |
| <b>PPO Hyperparameters (exploiters) (only differences to Main Agent hyperparameters)</b> |                                          |                   |                                                                                                       |
| PPO Exploiter                                                                            | exploiter_vf_coef                        | 0.1               | Weight of the value loss for the exploiters                                                           |

Some parameters (e. g. model paths, project names or logging names) were omitted for clarity.






Figure 10. Matchmaking process for LT. Dashed boxes are candidate opponent sets, while the percentages on the arrows are fixed sampling probabilities. The PFSP curve gives the weighting for the opponent win rate. We define  $w_{main\_hist}$  as the set of win rates of the Main Agent against all historical Main Exploiters.




Figure 11. Observation adjustment for player 2 in SP. To rotate the observations, the GridNet representation was flipped for each player-2 environment. Afterwards, the unit actions in the observation were flipped individually.

## References

- [Dhariwal et al., 2017] Dhariwal, P., Hesse, C., Klimov, O., Nichol, A., Plappert, M., Radford, A., Schulman, J., Sidor, S., Wu, Y., and Zhokhov, P. (2017). Openai baselines. <https://github.com/openai/baselines>. 3
- [Espeholt et al., 2018] Espeholt, L., Soyer, H., Munos, R., Simonyan, K., Mnih, V., Ward, T., Doron, Y., Firoiu, V., Harley, T., Dunning, I., Legg, S., and Kavukcuoglu, K. (2018). Impala: scalable distributed deep-rl with importance weighted actor-learner architectures. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, volume 80 of *Proceedings of Machine Learning Research*, pages 1406–1415. 3, 4
- [Goodfriend, 2024] Goodfriend, S. (2024). Raisocketai: The first deep reinforcement learning agent to win the iee microrts competition. In *2024 IEEE Conference on Games (CoG)*, pages 1–8. 3, 4, 5, 6, 10
- [Han et al., 2019] Han, L., Sun, P., Du, Y., Xiong, J., Wang, Q., Sun, X., Liu, H., and Zhang, T. (2019). Grid-wise control for multi-agent reinforcement learning in video game AI. In Chaudhuri, K. and Salakhutdinov, R., editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2576–2585. PMLR. 3, 5
- [Heinrich et al., 2015] Heinrich, J., Lanctot, M., and Silver, D. (2015). Fictitious self-play in extensive-form games. In Bach, F. and Blei, D., editors, *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pages 805–813. 6
- [Huang et al., 2021] Huang, S., Ontanón, S., Bamford, C., and Grela, L. (2021). Gym-microrts: Toward affordable full game real-time strategy games research with deep reinforcement learning. *CoRR*, abs/2105.13807. 1, 2, 3, 5, 6, 10
- [Huft, 2025] Huft, M. (2025). Game playing agents in microrts. 2, 3, 5, 9, 10, 11, 13, 15
- [Martin and Jhala, 2021] Martin, D. and Jhala, A. (2021). Beclone: Behavior cloning with inference for real-time strategy games. In *Joint Proceedings of the AIIDE 2021 Workshops*. 2, 3, 4, 9
- [Ontañón, 2013] Ontañón, S. (2013). The combinatorial multi-armed bandit problem and its application to real-time strategy games. In *Proceedings of the Ninth Conference on Artificial Intelligence and Interactive Digital Entertainment (AIIDE)*, volume 9, pages 58–64. AAAI. 3, 10
- [Schulman et al., 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms. *CoRR*, abs/1707.06347. 2, 9
- [Silver et al., 2018] Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., Lillicrap, T., Simonyan, K., and Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, 362:1140–1144. 6
- [Sutton and Barto, 1998] Sutton, R. S. and Barto, A. G. (1998). *Reinforcement Learning: An Introduction*. The MIT Press. 4
- [Torabi et al., 2018] Torabi, F., Warnell, G., and Stone, P. (2018). Behavioral cloning from observation. *arXiv preprint arXiv:1805.01954*. 2, 9
- [Vinyals et al., 2019] Vinyals, O., Babuschkin, I., Czarnecki, W. M., Mathieu, M., Dudzik, A., Chung, J., Choi, D. H., Powell, R., Ewalds, T., Georgiev, P., Oh, J., Horgan, D., Kroiss, M., Danihelka, I., Huang, A., Sifre, L., Cai, T., Agapiou, J. P., Jaderberg, M., Vezhnevets, A. S., Leblond, R., Pohlen, T., Dalibard, V., Budden, D., Sulsky, Y., Molloy, J., Paine, T. L., Gulcehre, C., Wang, Z., Pfaff, T., Wu, Y., Ring, R., Yogatama, D., Wünsch, D., McKinney, K., Smith, O., Schaul, T., Lillicrap, T., Kavukcuoglu, K., Hassabis, D., Apps, C., and Silver, D. (2019). Grandmaster level in starcraft ii using multi-agent reinforcement learning. *Nature*, 575(7782):350–354. 2, 4, 6, 7, 8, 9, 15, 16

## Erklärung

Diese Bachelor-Arbeit wurde teilweise unter Verwendung von ChatGPT zur Unterstützung bei Formulierungen, Rechtschreibung, Grammatik und Übersetzungen benutzt. Für die Implementierung eigener Ideen wurde im Code teilweise ChatGPT benutzt. Diese Inhalte wurden eigenständig geprüft und überarbeitet. Hiermit versichere ich, dass ich die Bachelor-Arbeit selbstständig verfasst und keine anderen als die angegebenen Quellen und Hilfsmittel benutzt habe.

Düsseldorf, den 01.04.2026

Niklas Glaser